Open ADKit Working Group Meeting 15/02/2024 #4182

oguzkaganozt started this conversation in **Working group meetings**

oguzkaganozt on Feb 15 Maintainer

Administrative

Open AD Kit Project Board

Attendees

Chaired by David Walmroth (PIX Moving) and Oğuz Öztürk (Leo Drive)

- [ ] Ryohsuke Mitsudome
- [ ] Kasper Mecklenburg
- [ ] Samet Kütük
- [ ] Christian John
- [ ] Esteve Fernandez
- [ ] Ambroise
- [ ] Bonolo
- [ ] Daniel Shish
- [ ] Hiroshi Igata
- [ ] Alex Carballo
- [ ] Heewon Kim
- [ ] Tanish

Agenda

- Introduction of new attendees
- **Open AD Kit 2024 Roadmap & Release Plan Updates** (Oguz, Samet, David, Kasper, ..)
  --Roadmap 2024: https://docs.google.com/spreadsheets/d/14y_STna--4EeNEEe_7htDxmfKjd927z6JaQPNcPVB44/edit#gid=0
  --Release Plan: https://docs.google.com/spreadsheets/u/3/d/12uZuwi0YRXVr987GTSER2-R7ejg9Igo91Fxt9Yk-ni8/edit?usp=drive_web&ouid=114941721738626529832
- **OAD Project Board Cleanup and Renaming WG tasks with clean descriptions** (Oguz, Kasper, ..)

Category

Working group meetings

Labels

meeting:openadkit-...

2 participants

- **Devops Dojos: Build/Run, container size optimization.** (Oguz) --
  https://github.com/orgs/autowarefoundation/discussions/3651
  -- Issue: Re-organize the Autoware Docker Containers #3999
  -- First PR: feat(docker): re-organize the autoware docker containers #4072 -- waiting for reviews to merge -- includes Dockerfile and Ansible modifications --
  -- Separated Autoware releases from Open AD Kit releases from github actions
  -- Tested building and running new containers with new build scripts runs fine
  -- Should we have perception artifacts in runtime containers ? Or should we make it optional ?

- **Devops Dojos: Ros Node Configuration Updates** (Kasper)
  -- Spreadsheet for node updates:
  https://docs.google.com/spreadsheets/d/1Cx3JgaSq28yzsUWsewWYOfC1fLs1UX0BiX-7HG89Ifw/edit#gid=1937796127
  -- DevOps Dashboard:
  https://github.com/orgs/autowarefoundation/projects/6/views/4

- **Proactively reaching out to individuals and organizations to participate in the WG** (Samet, David)
  -- Publishing newsletters or new media after the version release

- **Perception Containerization Discussion** (Oguz, David, Kasper, CJ ..)
  -- Initial approaches on how should perception stack be containerized
  -- Testing ORIN for perception running with x86 on PIX chassis (Mark, David)

- **AWSIM Integration and Testing** (Oguz, Mitsudome-san, Kasper)
  -- Validating new Open AD containers on the AWSIM environment
  -- Running AWSIM on x86 and Autoware on ARM

↑ 1

1 comment

Oldest | Newest | Top

doganulus on Feb 16 Collaborator edited

I have some follow-up questions on the perception container discussion for the perception team. Please let me know if there is another document to collect such questions but here they are:

- What is the minimum set of Nvidia dependencies for Cuda and TensorRT stack (assuming we want Nvidia GPU acceleration)? Could the perception team elaborate on runtime, build, and test dependencies?
- Could you elaborate on the role of `libnvinfer-plugin`, `libnvparsers`, `libnvonnxparsers` libraries on the perception system? Are they for runtime?
- Can there be an effort to refine the `findPackage(Cuda)`, `findPackage(CudaToolkit)`, `findPackage(TensorRT)` statements in

Cmake scripts? If the full installation is not needed, it is better to install those packages selectively and the build should not complain. In particular, the `tensorrt_common` package requires the full TensorRT setup as far as I can see.

- As we investigated Nvidia repos, I noticed relatively new `libnvinfer-lean` and `libnvinfer-dispatch` libraries, which are far smaller than regular TensorRT libraries. Can the perception team give its opinion on the use of these leaner runtimes and its potential performance implications if possible?

@oguzkaganozt **@mitsudome-r**

↑ 1 0 replies